

Recurrent Neural Networks for Language Modeling

Model modification

The plan is to modify the model build by adding L1 and L2 regularizations to the GRU layer and increase the number of RNN from 1024 to 2048. Reultes:

I noticed that the loss after 40th epoch the loss function was 3.

```
In [5]: import tensorflow as tf
        from tensorflow.keras.layers.experimental import preprocessing

        import numpy as np
        import os
        import time
```

```
In [6]: path_to_file = tf.keras.utils.get_file('shakespeare.txt', 'https://storage.googl

Downloading data from https://storage.googleapis.com/download.tensorflow.org/dat
a/shakespeare.txt
1122304/1115394 [=====] - 0s 0us/step
```

```
In [7]: path_to_file
```

```
Out[7]: '/root/.keras/datasets/shakespeare.txt'
```

```
In [8]: # Read, then decode for py2 compat.
        text = open(path_to_file, 'rb').read().decode(encoding='utf-8')
        # length of text is the number of characters in it
        print(f'Length of text: {len(text)} characters')
```

Length of text: 1115394 characters

```
In [9]: # print(text)
```

```
In [10]: # Take a look at the first 250 characters in text
         print(text[:250])
```

First Citizen:
Before we proceed any further, hear me speak.

All:
Speak, speak.

First Citizen:
You are all resolved rather to die than to famish?

All:

Resolved. resolved.

First Citizen:

First, you know Caius Marcius is chief enemy to the people.

```
In [11]: # The unique characters in the file
vocab = sorted(set(text))
print(f'{len(vocab)} unique characters')
```

65 unique characters

```
In [12]: print(vocab)
```

```
['\\n', ' ', '!', '$', '&', '"', '-', '3', ':', ';', '?', 'A', 'B',
'C', 'D', 'E', 'F', 'G', 'H', 'I', 'J', 'K', 'L', 'M', 'N', 'O', 'P', 'Q', 'R',
'S', 'T', 'U', 'V', 'W', 'X', 'Y', 'Z', 'a', 'b', 'c', 'd', 'e', 'f', 'g', 'h',
'i', 'j', 'k', 'l', 'm', 'n', 'o', 'p', 'q', 'r', 's', 't', 'u', 'v', 'w', 'x',
'y', 'z']
```

Process the text

```
In [13]: #The preprocessing.StringLookup layer can convert each character into
#a numeric ID. It just needs the text to be split into tokens first.
example_texts = ['abcdefg', 'xyz']

chars = tf.strings.unicode_split(example_texts, input_encoding='UTF-8')
chars
```

```
Out[13]: <tf.RaggedTensor [[b'a', b'b', b'c', b'd', b'e', b'f', b'g'], [b'x', b'y',
b'z']]>
```

```
In [14]: #Now create the preprocessing.StringLookup layer:
ids_from_chars = preprocessing.StringLookup(
    vocabulary=list(vocab))
```

```
In [15]: #It converts form tokens to character IDs, padding with 0
ids = ids_from_chars(chars)
ids
```

```
Out[15]: <tf.RaggedTensor [[41, 42, 43, 44, 45, 46, 47], [64, 65, 66]]>
```

```
In [16]: #invert this representation and recover human-readable strings
chars_from_ids = tf.keras.layers.experimental.preprocessing.StringLookup(
    vocabulary=ids_from_chars.get_vocabulary(), invert=True)
```

```
In [17]: #This layer recovers the characters from the vectors of IDs, and returns them as
chars = chars_from_ids(ids)
chars
```

```
Out[17]: <tf.RaggedTensor [[b'a', b'b', b'c', b'd', b'e', b'f', b'g'], [b'x', b'y',
b'z']]>
```

```
In [18]: #You can tf.strings.reduce_join to join the characters back into strings.
def text_from_ids(ids):
    return tf.strings.reduce_join(chars_from_ids(ids), axis=-1)
```

```
In [19]: text_from_ids(ids)
```

```
Out[19]: <tf.Tensor: shape=(2,), dtype=string, numpy=array([b'abcdefg', b'xyz'], dtype=object)>
```

The prediction task

```
In [20]: #Create training examples and targets
all_ids = ids_from_chars(tf.strings.unicode_split(text, 'UTF-8'))
all_ids
```

```
Out[20]: <tf.Tensor: shape=(1115394,), dtype=int64, numpy=array([20, 49, 58, ..., 47, 10, 2])>
```

```
In [21]: #use the tf.data.Dataset.from_tensor_slices function to convert the text vector
#a stream of character indices.
ids_dataset = tf.data.Dataset.from_tensor_slices(all_ids)
```

```
In [22]: for ids in ids_dataset.take(10):
    print(chars_from_ids(ids).numpy().decode('utf-8'))
```

F
i
r
s
t

C
i
t
i

```
In [23]: seq_length = 100
examples_per_epoch = len(text)//(seq_length+1)
```

```
In [24]: #batch method lets you easily convert these individual characters
#to sequences of the desired size.
sequences = ids_dataset.batch(seq_length+1, drop_remainder=True)

for seq in sequences.take(1):
    print(chars_from_ids(seq))
```

```
tf.Tensor(
[b'F' b'i' b'r' b's' b't' b' ' b'C' b'i' b't' b'i' b'z' b'e' b'n' b':'
 b'\n' b'B' b'e' b'f' b'o' b'r' b'e' b' ' b'w' b'e' b' ' b'p' b'r' b'o'
 b'c' b'e' b'e' b'd' b' ' b'a' b'n' b'y' b' ' b'f' b'u' b'r' b't' b'h'
 b'e' b'r' b',' b' ' b'h' b'e' b'a' b'r' b' ' b'm' b'e' b' ' b's' b'p'
 b'e' b'a' b'k' b'.' b'\n' b'\n' b'A' b'l' b'l' b':' b'\n' b'S' b'p' b'e'
 b'a' b'k' b',' b' ' b's' b'p' b'e' b'a' b'k' b'.' b'\n' b'\n' b'F' b'i'
```

```
b'r' b's' b't' b' ' b'C' b'i' b't' b'i' b'z' b'e' b'n' b':' b'\n' b'Y'
b'o' b'u' b' '], shape=(101,), dtype=string)
```

```
In [25]: #It's easier to see what this is doing if you join the tokens back into strings:
for seq in sequences.take(5):
    print(text_from_ids(seq).numpy())
```

```
b'First Citizen:\nBefore we proceed any further, hear me speak.\n\nAll:\nSpeak,
speak.\n\nFirst Citizen:\nYou '
b'are all resolved rather to die than to famish?\n\nAll:\nResolved. resolved.\n
\nFirst Citizen:\nFirst, you k'
b"now Caius Marcius is chief enemy to the people.\n\nAll:\nWe know't, we kno
w't.\n\nFirst Citizen:\nLet us ki"
b"ll him, and we'll have corn at our own price.\nIs't a verdict?\n\nAll:\nNo mor
e talking on't; let it be d"
b'one: away, away!\n\nSecond Citizen:\nOne word, good citizens.\n\nFirst Citize
n:\nWe are accounted poor citi'
```

```
In [26]: #function that takes a sequence as input, duplicates, and shifts it to align the
#label for each timestep
def split_input_target(sequence):
    input_text = sequence[:-1]
    target_text = sequence[1:]
    return input_text, target_text
```

```
In [27]: dataset = sequences.map(split_input_target)
```

```
In [28]: for input_example, target_example in dataset.take(1):
    print("Input :", text_from_ids(input_example).numpy())
    print("Target:", text_from_ids(target_example).numpy())
```

```
Input : b'First Citizen:\nBefore we proceed any further, hear me speak.\n\nAl
l:\nSpeak, speak.\n\nFirst Citizen:\nYou '
Target: b'irst Citizen:\nBefore we proceed any further, hear me speak.\n\nAll:\n
Speak, speak.\n\nFirst Citizen:\nYou '
```

Create training batches

```
In [29]: # Batch size
BATCH_SIZE = 64

# Buffer size to shuffle the dataset
# (TF data is designed to work with possibly infinite sequences,
# so it doesn't attempt to shuffle the entire sequence in memory. Instead,
# it maintains a buffer in which it shuffles elements).
BUFFER_SIZE = 10000

dataset = (
    dataset
    .shuffle(BUFFER_SIZE)
    .batch(BATCH_SIZE, drop_remainder=True)
    .prefetch(tf.data.experimental.AUTOTUNE))

dataset
```

Out[29]: <PrefetchDataset shapes: ((64, 100), (64, 100)), types: (tf.int64, tf.int64)>

Build The Model

```
In [30]: # Length of the vocabulary in chars
vocab_size = len(vocab)

# The embedding dimension
embedding_dim = 256

# Number of RNN units
rnn_units = 2048 #increase from 1024
```

```
In [31]: class MyModel(tf.keras.Model):
    def __init__(self, vocab_size, embedding_dim, rnn_units):
        super().__init__(self)
        self.embedding = tf.keras.layers.Embedding(vocab_size, embedding_dim)
        self.gru = tf.keras.layers.GRU(rnn_units,
                                         kernel_regularizer=tf.keras.regularizers.L1(0.01),
                                         activity_regularizer=tf.keras.regularizers.L2(0.01),
                                         return_sequences=True,
                                         return_state=True)
        self.dense = tf.keras.layers.Dense(vocab_size)

    def call(self, inputs, states=None, return_state=False, training=False):
        x = inputs
        x = self.embedding(x, training=training)
        if states is None:
            states = self.gru.get_initial_state(x)
        x, states = self.gru(x, initial_state=states, training=training)
        x = self.dense(x, training=training)

        if return_state:
            return x, states
        else:
            return x
```

```
In [32]: model = MyModel(
    # Be sure the vocabulary size matches the `StringLookup` layers.
    vocab_size=len(ids_from_chars.get_vocabulary()),
    embedding_dim=embedding_dim,
    rnn_units=rnn_units)
```

```
In [33]: for input_example_batch, target_example_batch in dataset.take(1):
    example_batch_predictions = model(input_example_batch)
    print(example_batch_predictions.shape, "# (batch_size, sequence_length, vocab_size)")

(64, 100, 67) # (batch_size, sequence_length, vocab_size)
```

```
In [34]: model.summary()

Model: "my_model"
```

Layer (type)	Output Shape	Param #
embedding (Embedding)	multiple	17152
gru (GRU)	multiple	14168064
dense (Dense)	multiple	137283
Total params: 14,322,499		
Trainable params: 14,322,499		
Non-trainable params: 0		

```
In [35]: sampled_indices = tf.random.categorical(example_batch_predictions[0], num_sampled=100)
        sampled_indices = tf.squeeze(sampled_indices, axis=-1).numpy()
```

```
In [36]: sampled_indices
```

```
Out[36]: array([52,  9,  2, 58,  0, 44,  2, 56,  6, 11,  0, 39, 29, 59, 56,  7, 60,
                52, 25, 18, 57, 16, 11, 28, 53, 64, 21, 54, 42,  2, 10, 44, 28, 34,
                35,  3, 33, 45, 59, 18, 49, 38, 53,  9, 53, 27,  5, 35, 15, 53, 62,
                39, 30, 23, 38, 57,  9, 33, 10, 36,  7, 54, 52, 12,  5, 39, 38, 64,
                52, 22, 57, 31,  1, 42, 63,  0, 25,  1,  2, 11, 60,  7, 51, 34, 30,
                1,  6,  7, 20, 35, 11, 33, 38, 12, 51,  8, 13, 57, 31, 32])
```

```
In [37]: print("Input:\n", text_from_ids(input_example_batch[0]).numpy())
        print()
        print("Next Char Predictions:\n", text_from_ids(sampled_indices).numpy())
```

Input:

b"th Baptista ta'en,\nThat none shall have access unto Bianca\nTill Katharina t
he curst have got a husba"

Next Char Predictions:

b"l-\nrd\np&3Y0sp'tlKDqB3NmxGnb\n.dNTU SesDiXm-mM\$UAmvYPIXq-S.V'nl:\$YXxlHqQ[UN
K]bwK[UNK]\n3t'kTP[UNK]&'FU3SX:k,;qQR"

Train the model

```
In [38]: #Attach an optimizer, and a loss function
        loss = tf.losses.SparseCategoricalCrossentropy(from_logits=True)
```

```
In [39]: example_batch_loss = loss(target_example_batch, example_batch_predictions)
        mean_loss = example_batch_loss.numpy().mean()
        print("Prediction shape: ", example_batch_predictions.shape, " # (batch_size, sequence_length, vocab_size)")
        print("Mean loss:      ", mean_loss)
```

```
Prediction shape: (64, 100, 67) # (batch_size, sequence_length, vocab_size)
Mean loss:      4.2048993
```

```
In [40]: tf.exp(mean_loss).numpy()
```

```
Out[40]: 67.01385
```

```
In [41]: #Configure the training procedure using the tf.keras.Model.compile method. Use t
#arguments and the loss function
model.compile(optimizer='adam', loss=loss)
```

Configure checkpoints

```
In [42]: #Use a tf.keras.callbacks.ModelCheckpoint to ensure that checkpoints are saved c
# Directory where the checkpoints will be saved
checkpoint_dir = './training_checkpoints'
# Name of the checkpoint files
checkpoint_prefix = os.path.join(checkpoint_dir, "ckpt_{epoch}")

checkpoint_callback = tf.keras.callbacks.ModelCheckpoint(
    filepath=checkpoint_prefix,
    save_weights_only=True)
```

Execute the training

```
In [43]: EPOCHS = 20
```

```
In [44]: history = model.fit(dataset, epochs=EPOCHS, callbacks=[checkpoint_callback])
```

```
Epoch 1/20
172/172 [=====] - 30s 159ms/step - loss: 4.6363
Epoch 2/20
172/172 [=====] - 29s 164ms/step - loss: 2.5367
Epoch 3/20
172/172 [=====] - 30s 168ms/step - loss: 2.2997
Epoch 4/20
172/172 [=====] - 31s 172ms/step - loss: 2.1452
Epoch 5/20
172/172 [=====] - 31s 176ms/step - loss: 2.0310
Epoch 6/20
172/172 [=====] - 32s 180ms/step - loss: 1.9454
Epoch 7/20
172/172 [=====] - 32s 182ms/step - loss: 1.8738
Epoch 8/20
172/172 [=====] - 32s 183ms/step - loss: 1.8186
Epoch 9/20
172/172 [=====] - 32s 183ms/step - loss: 1.7724
Epoch 10/20
172/172 [=====] - 32s 183ms/step - loss: 1.7284
Epoch 11/20
172/172 [=====] - 32s 183ms/step - loss: 1.6892
Epoch 12/20
172/172 [=====] - 32s 183ms/step - loss: 1.6509
Epoch 13/20
172/172 [=====] - 32s 183ms/step - loss: 1.6161
Epoch 14/20
172/172 [=====] - 32s 183ms/step - loss: 1.5840
Epoch 15/20
172/172 [=====] - 32s 183ms/step - loss: 1.5507
Epoch 16/20
172/172 [=====] - 33s 183ms/step - loss: 1.5194
Epoch 17/20
```

```

172/172 [=====] - 32s 183ms/step - loss: 1.4825
Epoch 18/20
172/172 [=====] - 32s 183ms/step - loss: 1.4470
Epoch 19/20
172/172 [=====] - 32s 183ms/step - loss: 1.4064
Epoch 20/20
172/172 [=====] - 32s 183ms/step - loss: 1.3621

```

Generate text

In []:

```

#The following makes a single step prediction:
class OneStep(tf.keras.Model):
    def __init__(self, model, chars_from_ids, ids_from_chars, temperature=1.0):
        super().__init__()
        self.temperature = temperature
        self.model = model
        self.chars_from_ids = chars_from_ids
        self.ids_from_chars = ids_from_chars

    # Create a mask to prevent "" or "[UNK]" from being generated.
    skip_ids = self.ids_from_chars(['', '[UNK]'][:, None])
    sparse_mask = tf.SparseTensor(
        # Put a -inf at each bad index.
        values=[-float('inf')]*len(skip_ids),
        indices=skip_ids,
        # Match the shape to the vocabulary
        dense_shape=[len(ids_from_chars.get_vocabulary())])
    self.prediction_mask = tf.sparse.to_dense(sparse_mask)

    @tf.function
    def generate_one_step(self, inputs, states=None):
        # Convert strings to token IDs.
        input_chars = tf.strings.unicode_split(inputs, 'UTF-8')
        input_ids = self.ids_from_chars(input_chars).to_tensor()

        # Run the model.
        # predicted_logits.shape is [batch, char, next_char_logits]
        predicted_logits, states = self.model(inputs=input_ids, states=states,
                                              return_state=True)

        # Only use the last prediction.
        predicted_logits = predicted_logits[:, -1, :]
        predicted_logits = predicted_logits/self.temperature
        # Apply the prediction mask: prevent "" or "[UNK]" from being generated.
        predicted_logits = predicted_logits + self.prediction_mask

        # Sample the output logits to generate token IDs.
        predicted_ids = tf.random.categorical(predicted_logits, num_samples=1)
        predicted_ids = tf.squeeze(predicted_ids, axis=-1)

        # Convert from token ids to characters
        predicted_chars = self.chars_from_ids(predicted_ids)

        # Return the characters and model state.
        return predicted_chars, states

```

In []:

```
one_step_model = OneStep(model, chars_from_ids, ids_from_chars)
```



```
In [ ]: start = time.time()
states = None
next_char = tf.constant(['ROMEO:'])
result = [next_char]

for n in range(1000):
    next_char, states = one_step_model.generate_one_step(next_char, states=states)
    result.append(next_char)

result = tf.strings.join(result)
end = time.time()
print(result[0].numpy().decode('utf-8'), '\n\n' + '_'*80)
print('\nRun time:', end - start)
```

ROMEO:
The last of Warwick call.

CAMILLO:
Nay, your office is dead.

LORTS:
So doth the death, they have a hearten mouth:
And how she'll sooner in his mother's.
The foul sir's news.

BIANCA:
Adgend and grief say he is colder,
For such a hand of her, of my heart
And hell in sucken me Above a cleagures
That you shall unto Lonnon o'er the land, whose hap that Clifford,
Was sentence of my life; for he dishollow'd attending music and
personal, sir.

GLOUCESTER:
Sir Richard, Or little tongue
For beying foot to such absolute.

LEONTES:
What willoub children how?

ROMEO:
I pray thee, my lord.

DUKE VINCENTIO:
It is not meet him well: you are dishonour'd between out
Where'er the people of this iclan careless.
Say that she lives.

TRASALLE:
I do become him; for I think, let me hear
The wind sid little eye o' the people, who haviness
my wag's faithful fearful long, and nothing else.'
Hast thou behold it straight degree?

HORTENSIO:
For this affection given him hence,
And he will muck in promise beaute

Run time: 1.754654884338379

In []:

